

LoongArch ELF ABI specification

Loongson Technology Corporation Limited

Version 2.01

Table of Contents

Register Convention	1
Aliases for return value registers	1
Type Size and Alignment	3
ELF Object Files	4
EI_CLASS: File class	4
e_machine: Identifies the machine	4
e_flags: Identifies ABI type and version	4
Relocations	6
Program Interpreter Path	13
Procedure Calling Convention	14
Abbreviations	14
Argument Registers	14
ABI LP64D	14
C Data Types and Alignment	14
Argument passing	14
Return values	17
Stack	17
Appendix: Revision History	18

Register Convention

Table 1. General-purpose Register Convention

Name	Alias	Meaning	Preserved across calls
<code>\$r0</code>	<code>\$zero</code>	Constant zero	(Constant)
<code>\$r1</code>	<code>\$ra</code>	Return address	No
<code>\$r2</code>	<code>\$tp</code>	Thread pointer	(Non-allocatable)
<code>\$r3</code>	<code>\$sp</code>	Stack pointer	Yes
<code>\$r4 - \$r5</code>	<code>\$a0 - \$a1</code>	Argument registers / return value registers	No
<code>\$r6 - \$r11</code>	<code>\$a2 - \$a7</code>	Argument registers	No
<code>\$r12 - \$r20</code>	<code>\$t0 - \$t8</code>	Temporary registers	No
<code>\$r21</code>		Reserved	(Non-allocatable)
<code>\$r22</code>	<code>\$fp / \$s9</code>	Frame pointer / Static register	Yes
<code>\$r23 - \$r31</code>	<code>\$s0 - \$s8</code>	Static registers	Yes

Table 2. Floating-point Register Convention

Name	Alias	Meaning	Preserved across calls
<code>\$f0 - \$f1</code>	<code>\$fa0 - \$fa1</code>	Argument registers / return value registers	No
<code>\$f2 - \$f7</code>	<code>\$fa2 - \$fa7</code>	Argument registers	No
<code>\$f8 - \$f23</code>	<code>\$ft0 - \$ft15</code>	Temporary registers	No
<code>\$f24 - \$f31</code>	<code>\$fs0 - \$fs7</code>	Static registers	Yes

Temporary registers are also known as caller-saved registers. Static registers are also known as callee-saved registers.

Aliases for return value registers

You may see the names `$v0 $v1 $fv0 $fv1` in some very early LoongArch assembly code; they simply alias to `$a0 $a1 $fa0 $fa1` respectively. The aliases are initially meant to match MIPS convention with separate argument / return value registers. However, because LoongArch actually has no dedicated return value registers, such usage may lead to confusion. Hence, it is not recommended to use these aliases.

Due to implementation details, it may not be easy to give a register multiple ABI names for a given downstream project. New programs processing LoongArch assembly should not support these aliases. Portable LoongArch assembly should avoid these aliases.

NOTE

For toolchain components provided by the Loongson Corporation, the migration procedure is:

Let the version of the component at this spec's effect date be N,

1. keep support in the version N and its stable branch,
2. warn on such usage in the version N+1,
3. remove support in the version N+2.

For the respective upstream projects of the components, the procedure above shall be followed if support for such usage is already upstream; "version N" shall be interpreted as the first release version containing LoongArch support in that case. For the components not yet upstream, and not interacting with other components that may expect such usage, support for such usage will never be implemented.

Type Size and Alignment

Table 3. LP64 Data Model (base ABI types: `lp64d` `lp64f` `lp64s`)

Scalar type	Size (Bytes)	Alignment (Bytes)
<code>bool</code> / <code>_Bool</code>	1	1
<code>unsigned char</code> / <code>char</code>	1	1
<code>unsigned short</code> / <code>short</code>	2	2
<code>unsigned int</code> / <code>int</code>	4	4
<code>unsigned long</code> / <code>long</code>	8	8
<code>unsigned long long</code> / <code>long long</code>	8	8
pointer types	8	8
<code>float</code>	4	4
<code>double</code>	8	8
<code>long double</code>	16	16

Table 4. ILP32 Data Model (base ABI types: `ilp32d` `ilp32f` `ilp32s`)

Scalar type	Size (Bytes)	Alignment (Bytes)
<code>bool</code> / <code>_Bool</code>	1	1
<code>unsigned char</code> / <code>char</code>	1	1
<code>unsigned short</code> / <code>short</code>	2	2
<code>unsigned int</code> / <code>int</code>	4	4
<code>unsigned long</code> / <code>long</code>	4	4
<code>unsigned long long</code> / <code>long long</code>	8	8
pointer types	4	4
<code>float</code>	4	4
<code>double</code>	8	8
<code>long double</code>	16	16

For all base ABI types of LoongArch, the `char` datatype is signed by default.

ELF Object Files

All common ELF definitions referenced in this section can be found in [the latest SysV gABI specification](#).

EI_CLASS: File class

Table 5. ELF file classes

EI_CLASS	Value	Description
ELFCLASS32	1	ELF32 object file
ELFCLASS64	2	ELF64 object file

e_machine: Identifies the machine

LoongArch (258)

e_flags: Identifies ABI type and version

Table 6. ABI-related bits in e_flags

Bit 31 - 8	Bit 7 - 6	Bit 5 - 3	Bit 2 - 0
(reserved)	ABI version	ABI extension	Base ABI Modifier

The ABI type of an ELF object is uniquely identified by EI_CLASS and e_flags[7:0] in its header.

Within this combination, EI_CLASS and e_flags[2:0] correspond to the **base ABI** type, where the expression of C integral and pointer types ([data model](#)) is uniquely determined by EI_CLASS value, and e_flags[2:0] represents additional properties of the base ABI type, including the FP calling convention. We refer to e_flags[2:0] as the **base ABI modifier**.

As a result, programs in lp64* / ilp32* ABI should only be encoded with ELF64 / ELF32 object files, respectively.

0x0 0x4 0x5 0x6 0x7 are reserved values for e_flags[2:0].

Table 7. Base ABI Types

Name	EI_CLASS	Base ABI Modifier (e_flags[2:0])	Description
lp64s	ELFCLASS64	0x1	Uses 64-bit GPRs and the stack for parameter passing. Data model is LP64, where long and pointers are 64-bit while int is 32-bit.
lp64f	ELFCLASS64	0x2	Uses 64-bit GPRs, 32-bit FPRs and the stack for parameter passing. Data model is LP64, where long and pointers are 64-bit while int is 32-bit.

Name	ELFCLASS	Base ABI Modifier (<code>e_flags[2:0]</code>)	Description
lp64d	ELFCLASS64	0x3	Uses 64-bit GPRs, 64-bit FPRs and the stack for parameter passing. Data model is LP64, where <code>long</code> and pointers are 64-bit while <code>int</code> is 32-bit.
ilp32s	ELFCLASS32	0x1	Uses 32-bit GPRs and the stack for parameter passing. Data model is ILP32, where <code>int</code> , <code>long</code> and pointers are 32-bit.
ilp32f	ELFCLASS32	0x2	Uses 32-bit GPRs, 32-bit FPRs and the stack for parameter passing. Data model is ILP32, where <code>int</code> , <code>long</code> and pointers are 32-bit.
ilp32d	ELFCLASS32	0x3	Uses 32-bit GPRs, 64-bit FPRs and the stack for parameter passing. Data model is ILP32, where <code>int</code> , <code>long</code> and pointers are 32-bit.

`e_flags[5:3]` correspond to the ABI extension type.

Table 8. ABI Extension types

Name	<code>e_flags[5:3]</code>	Description
base	0x0	No extra ABI features.
	0x1 - 0x7	(reserved)

`e_flags[7:6]` marks the ABI version of an ELF object.

Table 9. ABI Version

ABI version	Value	Description
v0	0x0	Stack operands base relocation type.
v1	0x1	Supporting relocation types directly writing to immediate slots. Can be implemented separately without compatibility with v0.
	0x2 0x3	Reserved.

Relocations

Table 10. ELF Relocation types

Enum	ELF reloc type	Usage	Detail
0	R_LARCH_NONE		
1	R_LARCH_32	Runtime address resolving	<code>*(int32_t *) PC = RtAddr + A</code>
2	R_LARCH_64	Runtime address resolving	<code>*(int64_t *) PC = RtAddr + A</code>
3	R_LARCH_RELATIVE	Runtime fixup for load-address	<code>*(void **) PC = B + A</code>
4	R_LARCH_COPY	Runtime memory copy in executable	<code>memcpy (PC, RtAddr, sizeof (sym))</code>
5	R_LARCH_JUMPSLOT	Runtime PLT supporting	<i>implementation-defined</i>
6	R_LARCH_TLS_DTPMOD32	Runtime relocation for TLS-GD	<code>*(int32_t *) PC = ID of module defining sym</code>
7	R_LARCH_TLS_DTPMOD64	Runtime relocation for TLS-GD	<code>*(int64_t *) PC = ID of module defining sym</code>
8	R_LARCH_TLS_DTPREL32	Runtime relocation for TLS-GD	<code>*(int32_t *) PC = DTV-relative offset for sym</code>
9	R_LARCH_TLS_DTPREL64	Runtime relocation for TLS-GD	<code>*(int64_t *) PC = DTV-relative offset for sym</code>
10	R_LARCH_TLS_TPREL32	Runtime relocation for TLE-IE	<code>*(int32_t *) PC = T</code>
11	R_LARCH_TLS_TPREL64	Runtime relocation for TLE-IE	<code>*(int64_t *) PC = T</code>
12	R_LARCH_IRELATIVE	Runtime local indirect function resolving	<code>*(void **) PC = (((void *) (*)()) (B + A)) ()</code>
... Reserved for dynamic linker.			
20	R_LARCH_MARK_LA	Mark la.abs	Load absolute address for static link.
21	R_LARCH_MARK_PCREL	Mark external label branch	Access PC relative address for static link.
22	R_LARCH_SOP_PUSH_PCREL	Push PC-relative offset	<code>push (S - PC + A)</code>
23	R_LARCH_SOP_PUSH_ABSOLUTE	Push constant or absolute address	<code>push (S + A)</code>
24	R_LARCH_SOP_PUSH_DUP	Duplicate stack top	<code>opr1 = pop (), push (opr1), push (opr1)</code>

Enum	ELF reloc type	Usage	Detail
25	R_LARCH_SOP_PUSH_GPREL	Push for access GOT entry	push (G)
26	R_LARCH_SOP_PUSH_TLS_TPREL	Push for TLS-LE	push (T)
27	R_LARCH_SOP_PUSH_TLS_GOT	Push for TLS-IE	push (IE)
28	R_LARCH_SOP_PUSH_TLS_GD	Push for TLS-GD	push (GD)
29	R_LARCH_SOP_PUSH_PLT_PCREL	Push for external function calling	push (PLT - PC)
30	R_LARCH_SOP_ASSERT	Assert stack top	assert (pop ())
31	R_LARCH_SOP_NOT	Stack top operation	push (!pop ())
32	R_LARCH_SOP_SUB	Stack top operation	opr2 = pop (), opr1 = pop (), push (opr1 - opr2)
33	R_LARCH_SOP_SL	Stack top operation	opr2 = pop (), opr1 = pop (), push (opr1 << opr2)
34	R_LARCH_SOP_SR	Stack top operation	opr2 = pop (), opr1 = pop (), push (opr1 >> opr2)
35	R_LARCH_SOP_ADD	Stack top operation	opr2 = pop (), opr1 = pop (), push (opr1 + opr2)
36	R_LARCH_SOP_AND	Stack top operation	opr2 = pop (), opr1 = pop (), push (opr1 & opr2)
37	R_LARCH_SOP_IF_ELSE	Stack top operation	opr3 = pop (), opr2 = pop (), opr1 = pop (), push (opr1 ? opr2 : opr3)
38	R_LARCH_SOP_POP_32_S_10_5	Instruction imm-field relocation	opr1 = pop (), (*(uint32_t *) PC) [14 ... 10] = opr1 [4 ... 0] with check 5-bit signed overflow
39	R_LARCH_SOP_POP_32_U_10_12	Instruction imm-field relocation	opr1 = pop (), (*(uint32_t *) PC) [21 ... 10] = opr1 [11 ... 0] with check 12-bit unsigned overflow

Enum	ELF reloc type	Usage	Detail
40	R_LARCH_SOP _POP_32_S_1 0_12	Instruction imm-field relocation	<pre>opr1 = pop (), (*(uint32_t *) PC) [21 ... 10] = opr1 [11 ... 0]</pre> <p>with check 12-bit signed overflow</p>
41	R_LARCH_SOP _POP_32_S_1 0_16	Instruction imm-field relocation	<pre>opr1 = pop (), (*(uint32_t *) PC) [25 ... 10] = opr1 [15 ... 0]</pre> <p>with check 16-bit signed overflow</p>
42	R_LARCH_SOP _POP_32_S_1 0_16_S2	Instruction imm-field relocation	<pre>opr1 = pop (), (*(uint32_t *) PC) [25 ... 10] = opr1 [17 ... 2]</pre> <p>with check 18-bit signed overflow and 4-bit aligned</p>
43	R_LARCH_SOP _POP_32_S_5 _20	Instruction imm-field relocation	<pre>opr1 = pop (), (*(uint32_t *) PC) [24 ... 5] = opr1 [19 ... 0]</pre> <p>with check 20-bit signed overflow</p>
44	R_LARCH_SOP _POP_32_S_0 _5_10_16_S2	Instruction imm-field relocation	<pre>opr1 = pop (), (*(uint32_t *) PC) [4 ... 0] = opr1 [22 ... 18],</pre> <pre>(*(uint32_t *) PC) [25 ... 10] = opr1 [17 ... 2]</pre> <p>with check 23-bit signed overflow and 4-bit aligned</p>
45	R_LARCH_SOP _POP_32_S_0 _10_10_16_S 2	Instruction imm-field relocation	<pre>opr1 = pop (), (*(uint32_t *) PC) [9 ... 0] = opr1 [27 ... 18],</pre> <pre>(*(uint32_t *) PC) [25 ... 10] = opr1 [17 ... 2]</pre> <p>with check 28-bit signed overflow and 4-bit aligned</p>
46	R_LARCH_SOP _POP_32_U	Instruction fixup	<pre>(*(uint32_t *) PC) = pop ()</pre> <p>with check 32-bit unsigned overflow</p>
47	R_LARCH_ADD 8	8-bit in-place addition	<pre>*(int8_t *) PC += S + A</pre>
48	R_LARCH_ADD 16	16-bit in-place addition	<pre>*(int16_t *) PC += S + A</pre>

Enum	ELF reloc type	Usage	Detail
49	R_LARCH_ADD 24	24-bit in-place addition	$*(\text{int24_t } *) \text{ PC } += \text{ S } + \text{ A}$
50	R_LARCH_ADD 32	32-bit in-place addition	$*(\text{int32_t } *) \text{ PC } += \text{ S } + \text{ A}$
51	R_LARCH_ADD 64	64-bit in-place addition	$*(\text{int64_t } *) \text{ PC } += \text{ S } + \text{ A}$
52	R_LARCH_SUB 8	8-bit in-place subtraction	$*(\text{int8_t } *) \text{ PC } -= \text{ S } + \text{ A}$
53	R_LARCH_SUB 16	16-bit in-place subtraction	$*(\text{int16_t } *) \text{ PC } -= \text{ S } + \text{ A}$
54	R_LARCH_SUB 24	24-bit in-place subtraction	$*(\text{int24_t } *) \text{ PC } -= \text{ S } + \text{ A}$
55	R_LARCH_SUB 32	32-bit in-place subtraction	$*(\text{int32_t } *) \text{ PC } -= \text{ S } + \text{ A}$
56	R_LARCH_SUB 64	64-bit in-place subtraction	$*(\text{int64_t } *) \text{ PC } -= \text{ S } + \text{ A}$
57	R_LARCH_GNU _VTINHERIT	GNU C++ vtable hierarchy	
58	R_LARCH_GNU _VTENTRY	GNU C++ vtable member usage	
... Reserved			
64	R_LARCH_B16	18-bit PC-relative jump	$*(\text{uint32_t } *) \text{ PC} [25 \dots 10] = (\text{S} + \text{A} - \text{PC}) [17 \dots 2]$ with check 18-bit signed overflow and 4-bit aligned
65	R_LARCH_B21	23-bit PC-relative jump	$*(\text{uint32_t } *) \text{ PC} [4 \dots 0] = (\text{S} + \text{A} - \text{PC}) [22 \dots 18],$ $*(\text{uint32_t } *) \text{ PC} [25 \dots 10] = (\text{S} + \text{A} - \text{PC}) [17 \dots 2]$ with check 23-bit signed overflow and 4-bit aligned
66	R_LARCH_B26	28-bit PC-relative jump	$*(\text{uint32_t } *) \text{ PC} [9 \dots 0] = (\text{S} + \text{A} - \text{PC}) [27 \dots 18],$ $*(\text{uint32_t } *) \text{ PC} [25 \dots 10] = (\text{S} + \text{A} - \text{PC}) [17 \dots 2]$ with check 28-bit signed overflow and 4-bit aligned
67	R_LARCH_ABS _HI20	[31 ... 12] bits of 32/64-bit absolute address	$*(\text{uint32_t } *) \text{ PC} [24 \dots 5] = (\text{S} + \text{A}) [31 \dots 12]$

Enum	ELF reloc type	Usage	Detail
68	R_LARCH_ABS_L012	[11 ... 0] bits of 32/64-bit absolute address	$(*(uint32_t *) PC) [21 \dots 10] = (S+A) [11 \dots 0]$
69	R_LARCH_ABS_64_L020	[51 ... 32] bits of 64-bit absolute address	$(*(uint32_t *) PC) [24 \dots 5] = (S+A) [51 \dots 32]$
70	R_LARCH_ABS_64_HI12	[63 ... 52] bits of 64-bit absolute address	$(*(uint32_t *) PC) [21 \dots 10] = (S+A) [63 \dots 52]$
71	R_LARCH_PCALA_HI20	[31 ... 12] bits of 32/64-bit PC-relative offset	$(*(uint32_t *) PC) [24 \dots 5] = (((S+A) \& \sim 0xffff) - (PC \& \sim 0xffff)) [31 \dots 12]$ Note: The lower 12 bits are not included when calculating the PC-relative offset.
72	R_LARCH_PCALA_L012	[11 ... 0] bits of 32/64-bit address	$(*(uint32_t *) PC) [21 \dots 10] = (S+A) [11 \dots 0]$
73	R_LARCH_PCALA64_L020	[51 ... 32] bits of 64-bit PC-relative offset	$(*(uint32_t *) PC) [24 \dots 5] = (S+A - (PC \& \sim 0xffffffff)) [51 \dots 32]$
74	R_LARCH_PCALA64_HI12	[63 ... 52] bits of 64-bit PC-relative offset	$(*(uint32_t *) PC) [21 \dots 10] = (S+A - (PC \& \sim 0xffffffff)) [63 \dots 52]$
75	R_LARCH_GOT_PC_HI20	[31 ... 12] bits of 32/64-bit PC-relative offset to GOT entry	$(*(uint32_t *) PC) [24 \dots 5] = (((GP+G) \& \sim 0xffff) - (PC \& \sim 0xffff)) [31 \dots 12]$
76	R_LARCH_GOT_PC_L012	[11 ... 0] bits of 32/64-bit GOT entry address	$(*(uint32_t *) PC) [21 \dots 10] = (GP+G) [11 \dots 0]$
77	R_LARCH_GOT_64_PC_L020	[51 ... 32] bits of 64-bit PC-relative offset to GOT entry	$(*(uint32_t *) PC) [24 \dots 5] = (GP+G - (PC \& \sim 0xffffffff)) [51 \dots 32]$
78	R_LARCH_GOT_64_PC_HI12	[63 ... 52] bits of 64-bit PC-relative offset to GOT entry	$(*(uint32_t *) PC) [21 \dots 10] = (GP+G - (PC \& \sim 0xffffffff)) [63 \dots 52]$
79	R_LARCH_GOT_HI20	[31 ... 12] bits of 32/64-bit GOT entry absolute address	$(*(uint32_t *) PC) [24 \dots 5] = (GP+G) [31 \dots 12]$
80	R_LARCH_GOT_L012	[11 ... 0] bits of 32/64-bit GOT entry absolute address	$(*(uint32_t *) PC) [21 \dots 10] = (GP+G) [11 \dots 0]$
81	R_LARCH_GOT_64_L020	[51 ... 32] bits of 64-bit GOT entry absolute address	$(*(uint32_t *) PC) [24 \dots 5] = (GP+G) [51 \dots 32]$
82	R_LARCH_GOT_64_HI12	[63 ... 52] bits of 64-bit GOT entry absolute address	$(*(uint32_t *) PC) [21 \dots 10] = (GP+G) [63 \dots 52]$

Enum	ELF reloc type	Usage	Detail
83	R_LARCH_TLS_LE_HI20	[31 ... 12] bits of TLS LE 32/64-bit offset from TP register	$(*(uint32_t *) PC) [24 \dots 5] = T [31 \dots 12]$
84	R_LARCH_TLS_LE_L012	[11 ... 0] bits of TLS LE 32/64-bit offset from TP register	$(*(uint32_t *) PC) [21 \dots 10] = T [11 \dots 0]$
85	R_LARCH_TLS_LE64_L020	[51 ... 32] bits of TLS LE 64-bit offset from TP register	$(*(uint32_t *) PC) [24 \dots 5] = T [51 \dots 32]$
86	R_LARCH_TLS_LE64_HI12	[63 ... 52] bits of TLS LE 64-bit offset from TP register	$(*(uint32_t *) PC) [21 \dots 10] = T [63 \dots 52]$
87	R_LARCH_TLS_IE_PC_HI20	[31 ... 12] bits of 32/64-bit PC-relative offset to TLS IE GOT entry	$(*(uint32_t *) PC) [24 \dots 5] = (((GP+IE) \& \sim 0xffff) - (PC \& \sim 0xffff)) [31 \dots 12]$
88	R_LARCH_TLS_IE_PC_L012	[11 ... 0] bits of 32/64-bit TLS IE GOT entry address	$(*(uint32_t *) PC) [21 \dots 10] = (GP+IE) [11 \dots 0]$
89	R_LARCH_TLS_IE64_PC_L020	[51 ... 32] bits of 64-bit PC-relative offset to TLS IE GOT entry	$(*(uint32_t *) PC) [24 \dots 5] = (GP+IE - (PC \& \sim 0xffffffff)) [51 \dots 32]$
90	R_LARCH_TLS_IE64_PC_HI12	[63 ... 52] bits of 64-bit PC-relative offset to TLS IE GOT entry	$(*(uint32_t *) PC) [21 \dots 10] = (GP+IE - (PC \& \sim 0xffffffff)) [63 \dots 52]$
91	R_LARCH_TLS_IE_HI20	[31 ... 12] bits of 32/64-bit TLS IE GOT entry absolute address	$(*(uint32_t *) PC) [24 \dots 5] = (GP+IE) [31 \dots 12]$
92	R_LARCH_TLS_IE_L012	[11 ... 0] bits of 32/64-bit TLS IE GOT entry absolute address	$(*(uint32_t *) PC) [21 \dots 10] = (GP+IE) [11 \dots 0]$
93	R_LARCH_TLS_IE64_L020	[51 ... 32] bits of 64-bit TLS IE GOT entry absolute address	$(*(uint32_t *) PC) [24 \dots 5] = (GP+IE) [51 \dots 32]$
94	R_LARCH_TLS_IE64_HI12	[63 ... 52] bits of 64-bit TLS IE GOT entry absolute address	$(*(uint32_t *) PC) [21 \dots 10] = (GP+IE) [63 \dots 52]$
95	R_LARCH_TLS_LD_PC_HI20	[31 ... 12] bits of 32/64-bit PC-relative offset to TLS LD GOT entry	$(*(uint32_t *) PC) [24 \dots 5] = (((GP+GD) \& \sim 0xffff) - (PC \& \sim 0xffff)) [31 \dots 12]$
96	R_LARCH_TLS_LD_HI20	[31 ... 12] bits of 32/64-bit TLS LD GOT entry absolute address	$(*(uint32_t *) PC) [24 \dots 5] = (GP+IE) [31 \dots 12]$
97	R_LARCH_TLS_GD_PC_HI20	[31 ... 12] bits of 32/64-bit PC-relative offset to TLS GD GOT entry	$(*(uint32_t *) PC) [24 \dots 5] = (((GP+GD) \& \sim 0xffff) - (PC \& \sim 0xffff)) [31 \dots 12]$
98	R_LARCH_TLS_GD_HI20	[31 ... 12] bits of 32/64-bit TLS GD GOT entry absolute address	$(*(uint32_t *) PC) [24 \dots 5] = (GP+IE) [31 \dots 12]$
99	R_LARCH_32_PCREL	32-bit PC relative	$(*(uint32_t *) PC) = (S+A-PC) [31 \dots 0]$

Enum	ELF reloc type	Usage	Detail
100	R_LARCH_REL AX	Instruction can be relaxed, paired with a normal relocation at the same address	

Program Interpreter Path

Table 11. Standard Program Interpreter Paths

Base ABI type	ABI extension type	Operating system / C library	Program interpreter path
lp64d	base	Linux, Glibc	/lib64/ld-linux-loongarch-lp64d.so.1
lp64f	base	Linux, Glibc	/lib64/ld-linux-loongarch-lp64f.so.1
lp64s	base	Linux, Glibc	/lib64/ld-linux-loongarch-lp64s.so.1
ilp32d	base	Linux, Glibc	/lib32/ld-linux-loongarch-ilp32d.so.1
ilp32f	base	Linux, Glibc	/lib32/ld-linux-loongarch-ilp32f.so.1
ilp32s	base	Linux, Glibc	/lib32/ld-linux-loongarch-ilp32s.so.1

Procedure Calling Convention

Abbreviations

In this document, **GRLEN** is the bit width of general-purpose register, **FRLEN** is the bit width of floating-point register and **WOA** is the bit width of the argument. The general-purpose argument register is denoted as **GAR** and the floating-point argument register is denoted as **FAR**.

Argument Registers

The basic principle of the LoongArch procedure calling convention is to pass arguments in registers as much as possible (i.e. floating-point arguments are passed in floating-point registers and non floating-point arguments are passed in general-purpose registers, as much as possible); arguments are passed on the stack only when no appropriate register is available.

The argument registers are:

1. Eight floating-point registers **fa0-fa7** used for passing pass floating-point arguments, and **fa0-fa1** are also used to return values.
2. Eight general-purpose registers **a0-a7** used for passing pass integer arguments, with **a0-a1** reused to return values.

Generally, the GARs are used to pass fixed-point arguments, and floating-point arguments when no FAR is available. Bit fields are stored in little endian. In addition, subroutines should ensure that the values of general-purpose registers **s0-s9** and floating-point registers **fs0-fs7** are preserved across procedure calls.

ABI LP64D

That is, **GRLEN** = 64, **FRLEN** = 64.

C Data Types and Alignment

The C data types and alignment in the LP64D ABI are defined in the [table 3](#).

In most cases, the unsigned integer data types are zero-extended when stored in general-purpose register, and the signed integer data types are sign-extended. However, in the **LP64D** ABI, unsigned 32-bit types, such as **unsigned int**, are stored in general-purpose registers as proper sign extensions of their 32-bit values.

Argument passing

Generally speaking, FARs are only used to pass floating-point arguments, GARs are used to pass non floating-point arguments and floating-point arguments when no FAR is available(**long double** type is also passed in a pair of GARs) and the reference.

Arguments passed by reference may be modified by the callee.

Scalar

There are two cases:

1. $0 < \text{WOA} \leq \text{GRLEN}$

- a. Argument is passed in a single argument register, or on the stack by value if none is available.
 - i. If the argument is floating-point type, the argument is passed in FAR. If no FAR is available, it's passed in GAR. If no GAR is available, it's passed on the stack. When passed in registers or on the stack, floating-point types narrower than GRLEN bits are widened to GRLEN bits, with the upper bits undefined.
 - ii. If the argument is integer or pointer type, the argument is passed in GAR. If no GAR is available, it's passed on the stack. When passed in registers or on the stack, the unsigned integer scalars narrower than GRLEN bits are zero-extended to GRLEN bits, and the signed integer scalars are sign-extended.

2. $GRLEN < WOA \leq 2 \times GRLEN$

- a. The argument is passed in a pair of GAR, with the low-order GRLEN bits in the lower-numbered register and the high-order GRLEN bits in the higher-numbered register. If exactly one register is available, the low-order GRLEN bits are passed in the register and the high-order GRLEN bits are passed on the stack. If no GAR is available, it's passed on the stack.

Structure

Empty structures are ignored by C compilers which support them as a non-standard extension (same as union arguments and return values). Bits unused due to padding, and bits past the end of a structure whose size in bits is not divisible by GRLEN, are undefined. And the layout of the structure on the stack is consistent with that in memory.

1. $0 < WOA \leq GRLEN$

- a. The structure has only fixed-point members. If there is an available GAR, the structure is passed through the GAR by value passing; If no GAR is available, it's passed on the stack.
- b. The structure has only floating-point members:
 - i. One floating-point member. The argument is passed in a FAR; If no FAR is available, the value is passed in a GAR; if no GAR is available, the value is passed on the stack.
 - ii. Two floating-point members. The argument is passed in a pair of available FAR, with the low-order **float** member bits in the lower-numbered FAR and the high-order **float** member bits in the higher-numbered FAR. If the number of available FAR is less than 2, it's passed in a GAR, and passed on the stack if no GAR is available.
- c. The structure has both fixed-point and floating-point members, i.e. the structure has one **float** member and...
 - i. Multiple fixed-point members. If there are available GAR, the structure is passed in a GAR, and passed on the stack if no GAR is available.
 - ii. Only one fixed-point member. If one FAR and one GAR are available, the floating-point member of the structure is passed in the FAR, and the integer member of the structure is passed in the GAR; If no floating-point register but one GAR is available, it's passed in GAR; If no GAR is available, it's passed on the stack.

2. $GRLEN < WOA \leq 2 \times GRLEN$

- a. Only fixed-point members.
 - i. The argument is passed in a pair of available GAR, with the low-order bits in the lower-numbered GAR and the high-order bits in the higher-numbered GAR. If only one GAR is available, the low-order bits are in the GAR and the high-order bits are on the stack, and passed on the stack if no GAR is available.
- b. Only floating-point members.
 - i. The structure has one **long double** member or one **double** member and two adjacent **float** members or 3-4 **float** members. The argument is passed in a pair of available GAR, with the low-order bits in the lower-numbered GAR and the high-order bits in the higher-numbered GAR. If only one GAR is available, the low-order bits are in the GAR and the high-order

bits are on the stack, and passed on the stack if no GAR is available.

- ii. The structure with two **double** members is passed in a pair of available FARs. If no a pair of available FARs, it's passed in GARs. A structure with one **double** member and one **float** member is same.
- c. Both fixed-point and floating-point members.
 - i. The structure has one **double** member and only one fixed-point member.
 - A. If one FAR and one GAR are available, the floating-point member of the structure is passed in the FAR, and the integer member of the structure is passed in the GAR; If no floating-point registers but two GARs are available, it's passed in the two GARs; If only one GAR is available, the low-order bits are in the GAR and the high-order bits are on the stack; And it's passed on the stack if no GAR is available.
 - ii. Others
 - A. The argument is passed in a pair of available GAR, with the low-order bits in the lower-numbered GAR and the high-order bits in the higher-numbered GAR. If only one GAR is available, the low-order bits are in the GAR and the high-order bits are on the stack, and passed on the stack if no GAR is available.

3. $WOA > 2 \times GRLEN$

- a. It's passed by reference and are replaced in the argument list with the address. If there is an available GAR, the reference is passed in the GAR, and passed on the stack if no GAR is available.

Structure and scalars passed on the stack are aligned to the greater of the type alignment and GRLEN bits, but never more than the stack alignment.

Union

Union is passed in GAR or stack.

1. $0 < WOA \leq GRLEN$

- a. The argument is passed in a GAR, or on the stack by value if no GAR is available.

2. $GRLEN < WOA \leq 2 \times GRLEN$

- a. The argument is passed in a pair of available GAR, with the low-order bits in the lower-numbered GAR and the high-order bits in the higher-numbered GAR. If only one GAR is available, the low-order bits are in the GAR and the high-order bits are on the stack. The arguments are passed on the stack when no GAR is available.

3. $WOA > 2 \times GRLEN$

- a. It's passed by reference and are replaced in the argument list with the address. If there is an available GAR, the reference is passed in the GAR, and passed on the stack if no GAR is available.

Complex

A complex floating-point number, or a structure containing just one complex floating-point number, is passed as though it were a structure containing two floating-point reals.

Variadic arguments

Variadic arguments are passed in GARs in the same manner as named arguments. And after a variadic argument has been passed on the stack, all future arguments will also be passed on the stack, i.e., the last argument register may be left unused due to the aligned register pair rule.

1. $0 < WOA \leq GRLEN$

- a. The variadic arguments are passed in a GAR, or on the stack by value if no GAR is available.

2. $\text{GRLEN} < \text{WOA} \leq 2 \times \text{GRLEN}$

- a. The variadic arguments are passed in a pair of GARs. If only one GAR is available, the low-order bits are in the GAR and the high-order bits are on the stack, and passed on the stack if no GAR is available. or on the stack by value if none is available. It should be noted that **long double** data type is passed in an aligned GAR pair (the first register in the pair is even-numbered).

3. $\text{WOA} > 2 \times \text{GRLEN}$

- a. It's passed by reference and are replaced in the argument list with the address. If there is an available GAR, the reference is passed in the GAR, and passed on the stack if no GAR is available.

Return values

1. Generally speaking, **a0** and **a1** are used to return non floating-point values, and **fa0** and **fa1** are used to return floating-point values.
2. Values are returned in the same manner as a first named argument of the same type would be passed. If such an argument would have been passed by reference, the caller allocates memory for the return value, and passes the address as an implicit first argument.
3. The reference of the return value is returned that is stored in GAR **a0** if the size of the return value is larger than $2 \times \text{GRLEN}$ bits.

Stack

1. In general, the stack frame for a subroutine may contain space to contain the following:
 - a. Space to store arguments passed to subroutines that this subroutine calls.
 - b. A place to store the subroutine's return address.
 - c. A place to store the values of saved registers.
 - d. A place for local data storage.
2. The stack grows downwards (towards lower addresses) and the stack pointer shall be aligned to a 128-bit boundary upon procedure entry. The first argument passed on the stack is located at offset zero of the stack pointer on function entry; following arguments are stored at correspondingly higher addresses.
3. Procedures must not rely upon the persistence of stack-allocated data whose addresses lies below the stack pointer.

Appendix: Revision History

- **v1.00**
 - Add register usage convention, data type conventions and the list of ELF relocation types.
- **v2.00**
 - Add description of ILP32 data model.
 - Add description of return value register aliases.
 - Add relocation types with direct immediate-filling semantics.
 - Add ABI version porting guidelines for toolchain implementations.
 - Add link to SysV gABI documentation.
 - Adjust asciidoc code style.
- **v2.01**
 - Adjust description of ABI type encoding scheme.
 - Add header for all tables.